

Homework 3

Problem 3.1

Solution:

I based my answers on this questions on these two websites: [sched-design-CFS.txt](#) and [Linux Journal](#)

- a) Fairness in the context of CFS means that the tasks get divided into categories and each group gets an equal part of the processor to process tasks.
- b) CFS scheduler selects tasks based on a such called $p \rightarrow se.vruntime$ value and it always tries to run the task with the smallest $p \rightarrow se.vruntime$ value.
The data structure chosen for this CFS is rbtree and it is chosen because it stores the tasks according to their $p \rightarrow se.vruntime$ value and when it needs to run a task it takes the most left one and runs it for some period of time until another task becomes the most left task.
- c) The CFS doesn't use time-slices and has no notion of them, although the previous scheduler did. The following parameters are affected by CFS time calculations:
 - There is a wait time for which each process has to wait until the CPU is assigned that task.
 - There is a *fair-clock* variable, which is basically what the CPU is defined to give to a certain process so that all processes have the same amount of processing. For example if there are 5 processes, each one would get 20% of the CPU power.
- d) Each task gets a weight according to its priority. So in that way the time would pass faster for low-priority tasks. This means that low-priority tasks get less CPU time rather than high-priority tasks.

Problem 3.2

Solution:

The code for this problem can be found in *runner.c*.